

13. Architectural Styles for Concurrency

Oscar Nierstrasz

Roadmap



- > What is Software Architecture?
- > Three-layered application architecture
- > Flow architectures
 - Active Prime Sieve
- > Blackboard architectures
 - Fibonacci with Linda

Sources

- > M. Shaw and D. Garlan, *Software Architecture: Perspectives on an Emerging Discipline*, Prentice-Hall, 1996.
- > F. Buschmann, et al., *Pattern-Oriented Software Architecture — A System of Patterns*, John Wiley, 1996.
- > D. Lea, *Concurrent Programming in Java — Design principles and Patterns*, The Java Series, Addison-Wesley, 1996.
- > N. Carriero and D. Gelernter, *How to Write Parallel Programs: a First Course*, MIT Press, Cambridge, 1990.

Roadmap



- > **What is Software Architecture?**
- > Three-layered application architecture
- > Flow architectures
 - Active Prime Sieve
- > Blackboard architectures
 - Fibonacci with Linda

Software Architecture

*A Software Architecture defines a system in terms of computational **components and interactions** amongst those components.*

*An Architectural Style defines a **family of systems** in terms of a pattern of structural organization.*

— cf. Shaw & Garlan, Software Architecture, pp. 3, 19

Architectural style

Architectural styles typically entail four kinds of properties:

- > A *vocabulary* of design elements
 - e.g., “pipes”, “filters”, “sources”, and “sinks”
- > A set of *configuration rules* that constrain compositions
 - e.g., pipes and filters must alternate in a linear sequence
- > A *semantic interpretation*
 - e.g., each filter reads bytes from its input stream and writes bytes to its output stream
- > A set of *analyses* that can be performed
 - e.g., if filters are “well-behaved”, no deadlock can occur, and all filters can progress in tandem

Roadmap





- > What is Software Architecture?
- > **Three-layered application architecture**
- > Flow architectures
 - Active Prime Sieve
- > Blackboard architectures
 - Fibonacci with Linda

Communication Styles

Shared Variables

Processes communicate indirectly.

Explicit synchronization mechanisms are needed.



Message-Passing

Communication and synchronization are combined.
Communication may be either *synchronous or asynchronous.*


Simulated Message-Passing

Most concurrency and communication styles can be simulated by one another:


Message-passing can be modeled by associating message queues to each process.

**Unsynchronized
objects**

**Synchronized
queues**



Three-layered Application Architectures



This kind of architecture avoids nested monitor problems by restricting concurrency control to a single layer.

Problems with Layered Designs

Hard to extend beyond three layers because:

- > Synchronization policies of different layers may conflict
 - E.g., nested monitor lockouts
- > Ground actions may need to know current policy
 - E.g., blocking vs. failing

Roadmap



- > What is Software Architecture?
- > Three-layered application architecture
- > **Flow architectures**
 - Active Prime Sieve
- > Blackboard architectures
 - Fibonacci with Linda

Flow Architectures

Many synchronization problems can be avoided by arranging things so that information only flows in one direction from sources to filters to sinks.

Unix “pipes and filters”:

- > Processes are connected in a linear sequence.

Control systems:

- > events are picked up by sensors, processed, and generate new events.

Workflow systems:

- > Electronic documents flow through workflow procedures.

Unix Pipes

Unix pipes are *bounded buffers* that connect producer and consumer processes (*sources, sinks and filters*):

```
cat file          # send file contents to output stream
| tr -c 'a-zA-Z' '\012' # put each word on one line
| sort            # sort the words
| uniq -c         # count occurrences of each word
| sort -rn        # sort in reverse numerical order
| more            # and display the result
```

Unix Pipes

Processes should *read from standard input* and *write to standard output* streams:

—Misbehaving processes give rise to “*broken pipes*”!

Process creation and *scheduling* are handled by the O/S.
Synchronization is handled implicitly by the I/O system
(through buffering).

Flow Stages

Every flow stage is a *producer* or *consumer* or both:

- > Splitters (Multiplexers) have *multiple successors*
 - Multicasters *clone results* to multiple consumers
 - Routers *distribute results* amongst consumers


- > Mergers (Demultiplexers) have *multiple predecessors*
 - Collectors *interleave inputs* to a single consumer
 - Combiners *process multiple input* to produce a single result

- > Conduits have both multiple predecessors and consumers

Flow Policies

Flow can be *pull-based*, *push-based*, or a mixture:

- > Pull-based flow: Consumers *take results* from Producers
- > Push-based flow: Producers *put results* to Consumers
- > Buffers:
 - Put-only buffers (relays) *connect push-based stages*
 - Take-only buffers (pre-fetch buffers) *connect pull-based stages*
 - Put-Take buffers connect (adapt) push-based stages to pull-based stages



Limiting Flow

Unbounded buffers:

- > If producers are faster than consumers, buffers may *exhaust available memory*

Unbounded threads:

- > Having too many threads can *exhaust system resources* more quickly than unbounded buffers

Bounded buffers:

- > Tend to be either *always full or always empty*, depending on relative speed of producers and consumers

Bounded thread pools:

- > *Harder to manage* than bounded buffers

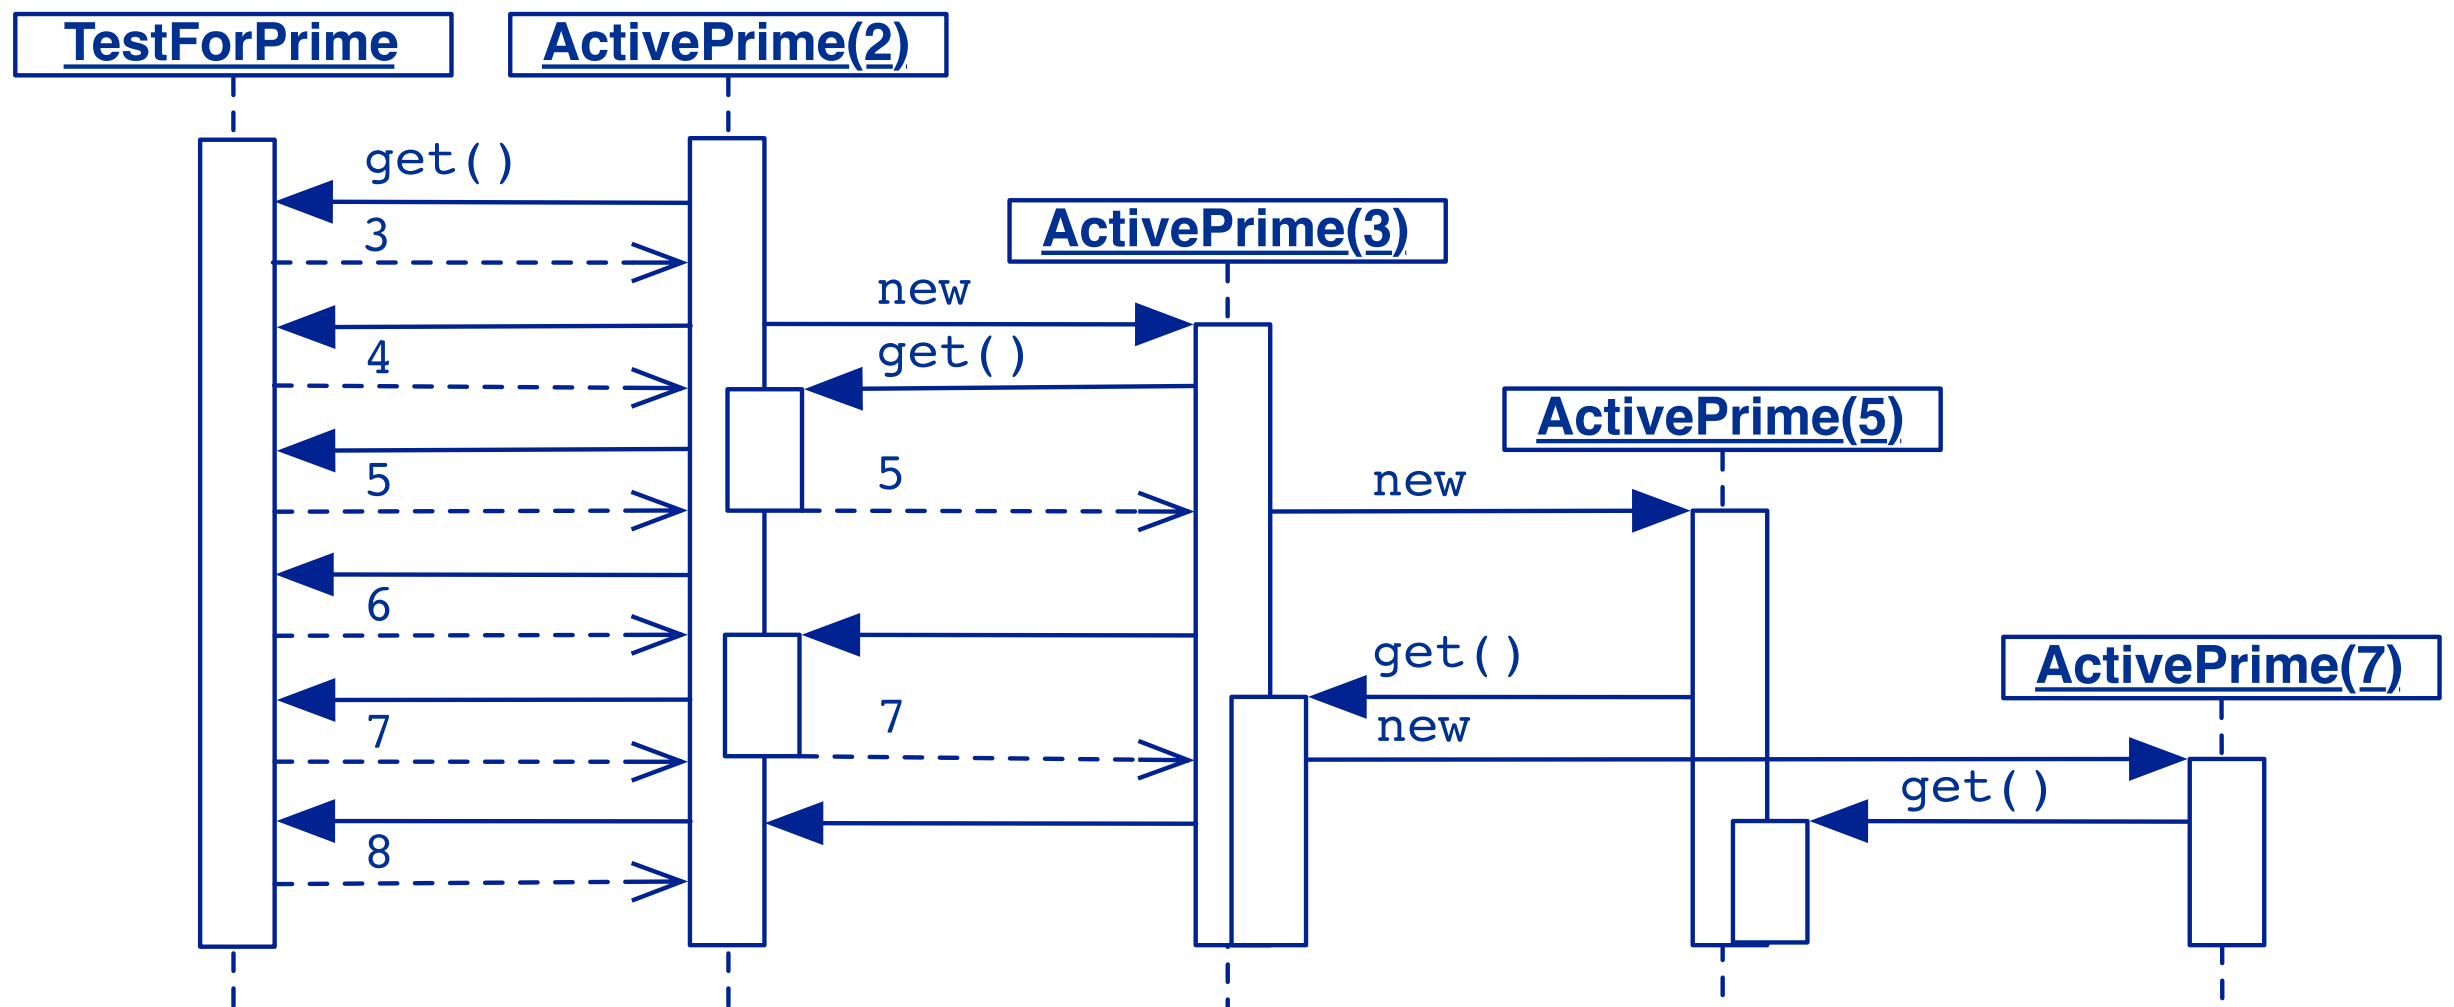
Roadmap



- > What is Software Architecture?
- > Three-layered application architecture
- > Flow architectures
 - **Active Prime Sieve**
- > Blackboard architectures
 - Fibonacci with Linda


Example: a Pull-based Prime Sieve

Primes are agents that reject non-primes, pass on candidates, or instantiate new prime agents:



Using Put-Take Buffers

Each ActivePrime uses a one-slot buffer to feed values to the next ActivePrime.



The first ActivePrime holds the seed value 2, gets values from a TestForPrime, and creates new ActivePrime instances whenever it detects a prime value.

The PrimeSieve

The main PrimeSieve class creates the initial configuration

```
public class PrimeSieve {  
    public static void main(String args[]) {  
        genPrimes(1000);  
    }  
    public static void genPrimes(int n) {  
        try {  
            ActivePrime firstPrime =  
                new ActivePrime(2, new TestForPrime(n));  
        } catch (Exception e) { }  
    }  
}
```

ActivePrimes



Pull-based integer sources

Active primes get values to test from an IntSource:

```
public interface Source<Value> { Value get(); }
class TestForPrime implements Source<Integer> {
    private int nextValue;
    private int maxValue;
    public TestForPrime(int max) {
        this.nextValue = 3;
        this.maxValue = max;
    }
    public Integer get() {
        if (nextValue < maxValue) { return nextValue++; }
        else { return 0; }
    }
}
```

The ActivePrime Class

ActivePrimes themselves implement IntSource

```
class ActivePrime extends Thread implements Source<Integer> {  
    private static Source<Integer> lastPrime; // shared  
    private int value; // value of this prime  
    private int square; // square of this prime  
    private Source<Integer> intSrc; // source of ints to test  
    private OneSlotBuffer<Integer> slot; // pass on test value  
    public ActivePrime(int value, Source<Integer> intSrc)  
        throws ActivePrimeFailure  
    {  
        this.value = value;  
        ...  
        slot = new OneSlotBuffer<Integer>();  
        lastPrime = this; // NB: set class variable  
        this.start();  
    }  
}
```



```
public int value() { return this.value; }
public void run() {
    int testValue = intSrc.get(); // may block
    while (testValue != 0) {
        if (testValue < this.square) {
            try {
                new ActivePrime(testValue, lastPrime);
            } catch (Exception e) {
                testValue = 0; // stop the thread
            }
        } else if ((testValue % this.value) > 0) {
            this.put(testValue);
        }
        testValue = intSrc.get(); // may block
    }
    put(0); // stop condition
}
private void put(Integer val) { slot.put(val); }
public Integer get() { return slot.get(); }
}
```


Roadmap



- > What is Software Architecture?
- > Three-layered application architecture
- > Flow architectures
 - Active Prime Sieve
- > **Blackboard architectures**
 - Fibonacci with Linda

Blackboard Architectures


Blackboard architectures put all synchronization in a “coordination medium” where agents can exchange messages.



Agents do not exchange messages directly, but post messages to the blackboard, and retrieve messages either by reading from a specific location (i.e., a channel), or by posing a query (i.e., a pattern to match).


Result Parallelism

Result parallelism is a blackboard architectural style in which workers *produce parts of a more complex whole*.



Agenda Parallelism


Agenda parallelism is a blackboard style in which workers *retrieve tasks to perform from a blackboard*, and may generate new tasks to perform.



Workers repeatedly retrieve tasks until everything is done.
Workers are typically able to perform arbitrary tasks.

Specialist Parallelism

Specialist parallelism is a style in which each worker is *specialized to perform a particular task*.



Specialist designs are *equivalent to message-passing*, and are often organized as *flow architectures*, with each specialist producing results for the next specialist to consume.

Linda

Linda is a *coordination medium*, with associated primitives for coordinating concurrent processes, that *can be added to an existing programming language*.

The coordination medium is a tuple-space, which can contain:

- data tuples* — tuples of primitives vales (numbers, strings ...)
- active tuples* — expressions which are evaluated and eventually turn into data tuples

Linda primitives

| | |
|-------------------------|---|
| <code>out (T)</code> | <i>output</i> a tuple T to the medium (non-blocking) e.g., <code>out("employee", "pingu", 35000)</code> |
| <code>in (S)</code> | <i>(destructively) input</i> a tuple matching S (blocking) e.g., <code>in("employee", "pingu", ?salary)</code> |
| <code>rd (S)</code> | <i>(non-destructively) read</i> a tuple (blocking) |
| <code>inp (S)</code> | <i>try to input</i> a tuple report success or failure (non-blocking) |
| <code>rdp (S)</code> | <i>try to read</i> a tuple report success or failure (non-blocking) |
| <code>eval (E)</code> | evaluate E in a <i>new process</i> leave the result in the tuple space |

Roadmap



- > What is Software Architecture?
- > Three-layered application architecture
- > Flow architectures
 - Active Prime Sieve
- > Blackboard architectures
 - Fibonacci with Linda**

Example: Fibonacci with JavaSpaces

```
BigInteger fib(final int n) {  
    Tuple tuple;  
    tuple = rdp(new Tuple("Fib", n, null));           // non-blocking  
    if (tuple != null) {  
        return tuple.result;  
    }  
    if (n < 2) {  
        out(new Tuple("Fib", n, BigInteger.ONE));    // non-blocking  
        return BigInteger.ONE;  
    }  
    eval("Fib", n, new Eval("fib(" + (n-1) + ") + fib(" + (n-2) + ")") {  
        public BigInteger expr() { return fib(n-1).add(fib(n-2)); }  
    } );  
    tuple = rd(new Tuple("Fib", n, null));           // blocking  
    return tuple.result;  
} // Post-condition: rdp("Fib",n,null) != null
```

JavaSpaces



Accessing the tuple space

```
public class Tuple {  
    public String functionName;  
    public Integer argument;  
    public BigInteger result;  
    ...  
}
```

```
private Tuple rdp(Tuple template) {  
    return tupleSpace.read(template, ZeroWaitTime);  
}  
private Tuple rd(Tuple template) {  
    return tupleSpace.read(template, WaitTime);  
}  
private Tuple inp(Tuple template) {  
    return tupleSpace.take(template, ZeroWaitTime);  
}  
private void out(Tuple template) {  
    tupleSpace.write(template, LeaseTime);  
}  
private void eval(String fn, final Integer arg, final Eval eval) {  
    new Thread() {  
        public void run() { out(new Tuple("Fib", arg, eval.expr())); }  
    }.start();  
}
```

We wrap a JavaSpaces implementation to resemble Linda more closely.

```
Computing fib(5)
rdp(Tuple("Fib", 5, null)) = null
eval("Fib", 5, fib(4)+fib(3))
rd(Tuple("Fib", 5, null)) [blocks]
rdp(Tuple("Fib", 4, null)) = null
eval("Fib", 4, fib(3)+fib(2))
rd(Tuple("Fib", 4, null)) [blocks]
rdp(Tuple("Fib", 3, null)) = null
eval("Fib", 3, fib(2)+fib(1))
rd(Tuple("Fib", 3, null)) [blocks]
rdp(Tuple("Fib", 2, null)) = null
eval("Fib", 2, fib(1)+fib(0))
rd(Tuple("Fib", 2, null)) [blocks]
rdp(Tuple("Fib", 1, null)) = null
out(Tuple("Fib", 1, 1))
rdp(Tuple("Fib", 0, null)) = null
out(Tuple("Fib", 0, 1))
out(Tuple("Fib", 2, 2))
rd(Tuple("Fib", 2, 2)) [returns]
rdp(Tuple("Fib", 1, null)) = Tuple("Fib", 1, 1)
out(Tuple("Fib", 3, 3))
rd(Tuple("Fib", 3, 3)) [returns]
rdp(Tuple("Fib", 2, null)) = Tuple("Fib", 2, 2)
out(Tuple("Fib", 4, 5))
rd(Tuple("Fib", 4, 5)) [returns]
rdp(Tuple("Fib", 3, null)) = Tuple("Fib", 3, 3)
out(Tuple("Fib", 5, 8))
rd(Tuple("Fib", 5, 8)) [returns]
DONE: fib(5) = 8
```

What you should know!

- > *What is a Software Architecture?*
- > *What are advantages and disadvantages of Layered Architectures?*
- > *What is a Flow Architecture? What are the options and tradeoffs?*
- > *What are Blackboard Architectures? What are the options and tradeoffs?*
- > *How does result parallelism differ from agenda parallelism?*
- > *How does Linda support coordination of concurrent agents?*

Can you answer these questions?

- > *How would you model message-passing agents in Java?*
- > *How would you classify Client/Server architectures?*
- > *Are there other useful styles we haven't yet discussed?*
- > *How can we prove that the Active Prime Sieve is correct?
Are you sure that new Active Primes will join the chain in the correct order?*
- > *Which Blackboard styles are better when we have multiple processors?*
- > *Which are better when we just have threads on a monoprocessor?*
- > *What will happen if you start two concurrent Fibonacci computations?*



Attribution-ShareAlike 3.0

You are free:

- to copy, distribute, display, and perform the work
- to make derivative works
- to make commercial use of the work

Under the following conditions:



Attribution. You must attribute the work in the manner specified by the author or licensor.



Share Alike. If you alter, transform, or build upon this work, you may distribute the resulting work only under a license identical to this one.

- For any reuse or distribution, you must make clear to others the license terms of this work.
- Any of these conditions can be waived if you get permission from the copyright holder.

Your fair use and other rights are in no way affected by the above.

<http://creativecommons.org/licenses/by-sa/3.0/>